

# CS232 Lab 3 Report

Part 3


Rijul Bhat (22B0971)

## Contents

<b>1</b>	<b>Measurements</b>	<b>2</b>
<b>2</b>	<b>Approach and Implementation</b>	<b>4</b>
<b>3</b>	<b>Observations and Conclusion</b>	<b>6</b>

# 1 Measurements

Plot: Time v/s N



TSC = 1002.460571 MHz.

Accessing Row Wise

N	no. of cycles	time = no of cycles/TSC(s)
128	299141	0.00029840674900718866
256	980566	0.0009781591699131278
512	3593883	0.003585061701144952
1024	14606570	0.01457071771454241
2048	56135762	0.05599797500663994

Table 1: Calculating using Row Wise Access

Accessing Column Wise

<b>N</b>	<b>no. of cycles</b>	<b>time = no of cycles/TSC(s)</b>
128	328575	0.0003277685023284572
256	1383500	0.001380104155737413
512	6868916	0.006852056029643085
1024	76376583	0.0761891142748995
2048	369593316	0.36868613758176433

Table 2: Calculating using Column Wise Access

## 2 Approach and Implementation

- **Overview**

The difference between the approaches in row and column execution is that accessing the indices is much easier in case of rows. In case of rows, we just need to iterate through them in the order in which they are stored. In case of column wise iteration, the data being stored in row major format each time we have to iterate we have to access the indices in increment intervals of  $8n$  bytes where  $nxn$  is the dimension of matrix and then come back to the first row every a column has been iterated which is inefficient and cumbersome.

Also, another important observation was that r12 has been used in the testbench for calculating number of cycles so if we are using r12 in our program we must push and pop and in the stack to preserve its value for the testbench.

- **Linear Combination**

Registers r14 and r15 are being used as the 2 loop counters i and j, the increment value i.e.  $ni+j$  (in case of row wise iteration) is stored in rbx. [a1], [a2] are stored in r13, r12 and [b1] is stored in r10, incrementing those addresses by rbx to give access to the corresponding

```

1 push r12
2 mov r14, 0
3
4 counter_i:
5 cmp r14, r9
6 je exit_i
7     mov r15, 0
8
9     mov rbx, r14
10    imul rbx, r9
11    sal rbx, 3
12
13    counter_j:
14    cmp r15, r9
15    je exit_j
16
17    mov r13, rdi
18    mov r12, rdx
19    mov r10, r8
20
21    add r13, rbx
22    add r12, rbx
23    add r10, rbx
24
25    mov r13, [r13]
26    mov r12, [r12]
27
28    imul r13, rsi
29    imul r12, rcx
30    add r13, r12
31
32    mov [r10], r13
33
34    add rbx, 8
35
36    add r15, 1
37    jmp counter_j
38
39 exit_j:
40 add r14, 1
41 jmp counter_i
42 exit_i:
43 pop r12

```

```

1 push r12
2 mov r14, 0
3
4 counter_i:
5 cmp r14, r9
6 je exit_i
7     mov r15, 0
8
9     mov rbx, r14
10    sal rbx, 3
11
12    counter_j:
13
14    cmp r15, r9
15    je exit_j
16
17    mov r13, rdi
18    mov r12, rdx
19    mov r10, r8
20
21    add r13, rbx
22    add r12, rbx
23    add r10, rbx
24
25    mov r13, [r13]
26    mov r12, [r12]
27
28    imul r13, rsi
29    imul r12, rcx
30    add r13, r12
31
32    mov [r10], r13
33
34    shl r9, 3
35    add rbx, r9
36    shr r9, 3
37    add r15, 1
38    jmp counter_j
39
40 exit_j:
41 add r14, 1
42 jmp counter_i
43 exit_i:
44 pop r12

```

index. Now value stored in r12 and r13 is multiplied by the scalars rcx, rsi and added then

stored in [r10] i.e. address of index i and j for [b1].

### • Hadamard Product

The program is extremely similar to the previous part so only the program inside the for loop is differing, here the product of r12 and r13 is to be taken instead of linear combination and store it in r10.

<pre> 1      mov r13, rdi 2      mov r12, rdx 3      mov r10, r8 4 5      add r13, rbx 6      add r12, rbx 7      add r10, rbx 8 9      mov r13, [r13] 10     mov r12, [r12] 11 12     imul r13, r12 13 14     mov [r10], r13 15 16     add rbx, 8 </pre>	<pre> 1      mov r13, rdi 2      mov r12, rdx 3      mov r10, r8 4 5      add r13, rbx 6      add r12, rbx 7      add r10, rbx 8 9      mov r13, [r13] 10     mov r12, [r12] 11 12     imul r13, r12 13 14     mov [r10], r13 15     shl r9, 3 16     add rbx, r9 17     shr r9, 3 </pre>
---	---

### • Alternate Sum

Here I have maintained value of i+j in rcx and exploited the fact that the least significant bits of odd decimal number is 1 and even is 0, so we can just **and** the value in rcx by 1, and determine whether we have to multiply by -1 or not while taking the sum.

<pre> 1      mov r13, rdi 2      add r13, rbx 3      mov r13, [r13] 4 5      mov rcx, r15 6      add rcx, r14 7 8      and rcx, 1 9 10     cmp rcx, 0 11     jne one 12 13     add rax, r13 14 15     jmp skip_one 16 17     one: 18 19     imul r13, -1 20     add rax, r13 21 22     skip_one: 23 24     add rbx, 8 </pre>	<pre> 1      mov r13, rdi 2      add r13, rbx 3      mov r13, [r13] 4 5      mov rcx, r15 6      add rcx, r14 7 8      and rcx, 1 9 10     cmp rcx, 0 11     jne one 12 13     add rax, r13 14 15     jmp skip_one 16 17     one: 18 19     imul r13, -1 20     add rax, r13 21 22     skip_one: 23     shl r9, 3 24     add rbx, r9 25     shr r9, 3 </pre>
--	--

### • Time, TSC and Cycle Count

```

1  #include <stdio.h>
2  #include <stdint.h>
3  #include <unistd.h>
4
5  static inline uint64_t rdtsc(void)
6  {

```

```
7 uint32_t lo,hi;
8 __asm__ __volatile__ ("rdtsc": "=a" (lo), "=d" (hi));
9 return ((uint64_t)hi << 32) | lo;
10 }
11
12 int main(void)
13 {
14     uint64_t tsc1 = rdtsc();
15     sleep(1);
16     uint64_t tsc2 = rdtsc();
17     printf("%f\n", ((float) (tsc2 - tsc1))/(1000000));
18     return 0;
19 }
```

The above C code computes the clock frequency of our processor using the TSC register and returns the output in MHz. I got this value close to **1000 MHz** every time the given program was executed on my Macbook Air using Docker.

### 3 Observations and Conclusion

We conclude that the row-wise access performs significantly better than the column-wise access. Row-wise access patterns take advantage of spatial proximity and cache line prefetching, resulting in fewer cache misses and quicker data retrieval, making them ideal for cache memory behaviour. In contrast, because column-wise access skips between memory regions and only partially utilizes spatial locality, it frequently results in more cache misses and is slower.